

QUERY EXECUTION

We want to translate the SQL query into an algebraic query to define the order of operations.

Then we do a logical transformation, typically push the selection operation towards the leaves, so that the join is computed later.

Finally, you generate a feasible access plan, choosing the correct algorithm, meaning that we annotate each node with the algo that we are going to use.

The order of the join is the most important decision in terms of speed of the query.



Physical operators

Each physical operator is a specific implementation of a logical operator.

- $R \rightarrow \text{TableScan}(R)$
- $\pi_x(E) \rightarrow \text{Project}(O_E, X) \rightarrow \text{argument } O_E \text{ and parameter } X$
- $E_1 \bowtie_{\text{cond}} E_2 \rightarrow \text{Nested Loop}(O_1, O_2, \text{cond})$
 - $\rightarrow \text{Index Nested Loop}(O_1, O_2, \text{cond})$
 - $\rightarrow \text{MergeJoin}(O_1, O_2, \text{cond})$

...

How do you implement a tree of operators?

One way is "one operator at a time"
or "materialization of the result of each operator."
Very simple to implement, but requires a lot of read/write.
Not how a DBMS works.

What is used is the iterator open/next/close,
called "iterator style" or "cursor interface".

For every node we use the iterator, for example

$\text{filter}(\text{Province} = P1) \rightarrow$ it just calls next on the TableScan,
which remembers the last cursor position
|
TableScan(student)

The filter doesn't need to keep everything in memory.
The TableScan just stores the position of the cursor.

We write the access plan like a tree,
where the argument of an operator is its child.

Children of operators

Observe that each operator has
a fixed number of children.

For example

- TableScan is always a leaf.

- Index Filter (R, Ix, ψ) \leadsto it gets out of the index Ix
all the values that satisfy
the condition ψ

\leadsto Index Filter (Students, IdAgeStud, age > 25)

open the index IdAgeStud,
retrieve RIDs with age > 25,
reads corresponding students from
the table Students

\leadsto this is ALWAYS a leaf,
because it always reads from disk,
cannot be combined a priori

Project and Filter

They are the simplest operators.

$\pi_A^b \rightarrow \text{Project}(\{A\})$

$\sigma_\psi \rightarrow \text{Filter}(\psi)$

$R \rightarrow \text{TableScan}(R)$

↑ Logical Plan

↑ Access Plan

π_A^b because SQL produces mult. sets
(no duplicate elimination)

~ this is a basic
SELECT
FROM
WHERE

Distinct elimination

If we want to do duplicate elimination, it is more complicated

$\pi_A \rightarrow \begin{cases} \text{Distinct} \\ \text{Sort}(\{A\}) \\ \text{Project}(\{A\}) \end{cases}$

$\sigma_\psi \rightarrow \text{Filter}(\psi)$

$R \rightarrow \text{TableScan}(R)$

This is needed because whenever you ask 'next' it remembers only the last element, so distinct can compare only with the last element

What is the cost of distinct? \emptyset because it happens in the main memory.

And Sort? It is $2 \cdot N \log N$, because it needs the disk,

so we try to reduce the use of sort,

and avoid it if data is already sorted

And Project? \emptyset

And Filter? \emptyset

And TableScan? $N \log N$

Selection and Index Filter

SELECT *
FROM R
WHERE A BETWEEN 50 AND 100

→ with idx on index on A

$\pi_x^b \rightarrow \text{Project} (*)$
 $\left. \begin{array}{c} | \\ \sigma_\psi \\ | \\ R \end{array} \right\} \rightarrow \text{IndexFilter}(R, \text{idx}, \psi)$

'IndexOnlyFilter' combines access to index with projection.

It just gives the values that it finds in the index, without going to the data. Much more efficient.

Go to the index and project to the same attributes of the index.
It does not read the table.

What is the difference between $\text{Filter}(O, \psi)$ and $\text{IndexFilter}(R, \text{idx}, \psi)$?

The first filter the records on the given argument O , while the second access the records of R through the index, with no argument.

all depends on the selectivity factor

About the cost,

$\text{Filter}(\text{TableScan}(R), \psi) \rightsquigarrow N_{\text{pag}}(R)$ ✓

$\text{IndexFilter}(R, \text{idx}, \psi) \rightsquigarrow C_I + C_D = C_I + sf * N_{\text{Rec}}(R)$

In addition, $\text{IndexFilter}(R, \text{idx}, \psi)$ gives data already sorted by A , then it might be not need to sort again.

Group By

The canonical way is by sorting and then break when the attribute changes.

```
SELECT A, COUNT(*)  
FROM R  
WHERE  
GROUP BY
```

$\sigma_{\text{COUNT}(*)>1} \rightarrow \text{Filter}(\text{COUNT}(*), 1)$
|

$\{A\} \times \{\text{COUNT}(*)\} \rightarrow \left\{ \begin{array}{l} \text{GroupBy}(\{A\}, \text{COUNT}(*)) \\ \text{Sort} \end{array} \right.$
|

$\sigma_{\psi} \rightarrow \text{Filter}(\psi)$
|

$R \rightarrow \text{Table Scan}(R)$

GROUPED means that when there is a break I will never get that value again.

SORTING implies GROUPED, it is stronger.

If you are not sure that data is grouped, then you need to sort.

But since grouped data is enough, it might not be required to sort.

Sort

The sort operator is $\tau_{\{A_i\}}$
and its cost is $2 * N \log$

Nested Loop (JOIN)

This is the simplest and most expensive algo.

Algo:

for each record r in R do \leftarrow outer relation
 for each record in S do \leftarrow inner relation
 if $r_i = s_j$ then add $\langle r, s \rangle$ to the result

The cost is $N_{\text{pag}}(R) + N_{\text{Rec}}(R) * N_{\text{Pag}}(S)$

\uparrow rest left
side

\uparrow for every record
on the left side
(outer rel.)

\uparrow we scan the entire
inner relation


\therefore the cost is quadratic, and you will never use it.

Page Nested Loop (JOIN)

Instead of scanning the entire relation once for every record of the outer relation, you read in entire page of the outer relation and join the entire page.

More difficult to implement and still quadratic,
but less expensive since $N_{\text{page}} < N_{\text{recs}}$.

The cost is $N_{\text{pag}}(R) + N_{\text{Pag}}(R) * N_{\text{Pag}}(S)$



Index Nested Loop

for every tuple of R , read the tuple $R.A$,
then use the index of $S.B$.

Index Nested Loop requires two arguments,
on the left, something that allows "next", like TableScan,
on the right, $\text{IndexFilter}(S, \text{Idx}, \underline{S.B = R.A})$ ←

This is the only operator where
data flows also from left to right,
not not only bottom-up.
It is the 'open' operator that carries
this information.

observe that here $S.B$
opens the index, is a constant,
while $R.A$ is the attribute

→ remember that for every
operator we do
open | next | close


→ the condition is

77' and before

Few minutes missing here

Merge Join (Join 4) (also Soft Merge)

If both tables have many records.
The cost is just that of sort,
that is $2 * N_{pg}$, linear in the
number of pages.



! Before MJ, be sure that
the input is sorted
exactly on the join attribute.

Should we use Index Nested Loop or Merge join?

The cost is respectively $N_{pg}(R) + k * N_{rec}(R)$ and $2 * N_{pg}$.

So if S has $N_{rec} = 100.000$ and $N_{pg} = 1.000$
and R has $N_{rec} = 3$ and $N_{pg} = 1$.

Then INL is very convenient, because the cost is $1 + k * 3$.

While MJ's cost is $2 * 1.000$.

But if R is as big as S, like

R with $N_{rec} = 100.000$ and $N_{pg} = 1.000$,

Then INL will be around 200.000

and MJ is $2 * 1.000 = 2.000$

In the transactional application, using few records, use INL,
while analytical applications use MJ.

Of course, to use INL, we need an INL.

But typically every DBMS put an index on every key and foreign key.

Operations of JOIN (review)

1. Nested Loop (O_E, O_I, Ψ)
 2. PsyeNested Loop (O_E, O_I, Ψ)
 3. Index Nested Loop (O_E, O_I, Ψ)
 4. Sort Merge (O_E, O_I, Ψ)
- } NOT USED
- } USED

Query Optimization

Typically means moving the restriction σ before the join, that is expensive. This is the logical optimization.

Then we have the physical plan generation of the plan with the optimal cost.

Actually we just avoid terrible plans.